

A dark blue vertical bar runs down the left side of the page. A blue arrow points to the right from the bar, containing the year 2024.

**2024**

# **Sentiment Analysis on News Related to Apple**

Capstone Project - Natural  
Language Processing

Several thin, curved lines in shades of blue and grey originate from the bottom left corner and sweep upwards and to the right.

**DIBOSH BARUAH**

## 1. Introduction

Sentiment analysis is the process of determining whether a piece of text expresses positive, negative, or neutral sentiment. In the context of financial markets, sentiment analysis on news articles related to companies can provide valuable insights into how news events impact stock prices and investor perceptions. This project aims to analyze news headlines and summaries related to Apple Inc. (AAPL) using sentiment analysis, and explore how the sentiment may correlate with stock market movements. The goal is to apply Natural Language Processing (NLP) techniques to assess the polarity (positive or negative sentiment) of news content and investigate whether sentiment influences stock price changes.

The project uses the **Finnhub API** to fetch news data for Apple Inc. and the **TextBlob** library for sentiment analysis. By applying NLP techniques, this project attempts to understand the sentiment expressed in news headlines and summaries and analyze its potential impact on the stock market price of Apple.

## 2. Data Collection

The data for this project is collected through the **Finnhub API**, which provides access to real-time news data related to various companies. Specifically, the news articles related to Apple Inc. (AAPL) were fetched using the `company_news` function from the Finnhub API. The data consists of the following attributes:

- **Headline:** The title of the news article.
- **Summary:** A brief description or summary of the article.
- **Datetime:** The publication date and time of the news article.
- **Source:** The news outlet or website from which the article originates.

The data collected covers a time range from **January 1, 2024, to January 31, 2024**, allowing for a focused analysis of sentiment over a one-month period.

### 3. Data Preprocessing

Several preprocessing steps were applied to the news data to make it ready for sentiment analysis:

#### 3.1 Handling Missing Values

The dataset was checked for missing values, particularly in the "headline" and "summary" columns. Any rows with missing values were removed to ensure a clean dataset for analysis.

#### 3.2 Text Cleaning

The news headlines and summaries were cleaned by removing any irrelevant characters such as special symbols, extra spaces, and stop words. This step ensures that the text is standardized and ready for sentiment analysis.

#### 3.3 Sentiment Analysis

The main feature of this project is performing sentiment analysis on the text data. The **TextBlob** library was used to calculate the **polarity score** of the headlines and summaries. The polarity score ranges from -1 (negative sentiment) to +1 (positive sentiment), with values close to 0 indicating neutral sentiment.

The sentiment analysis is applied as follows:

- **Headline Sentiment:** Sentiment score calculated for the headline of each news article.
- **Summary Sentiment:** Sentiment score calculated for the summary of each news article.

### 4. Model Building

The goal of this project is to investigate the relationship between sentiment and stock market movements. To achieve this, the following steps were taken:

#### 4.1 Stock Price Data Collection

Alongside the news data, Apple Inc.'s stock price data (AAPL) was also fetched using the **Finnhub API**. The stock data includes:

- **Date:** The date on which the stock data was recorded.
- **Open:** The opening price of AAPL stock on that date.
- **Close:** The closing price of AAPL stock on that date.
- **High:** The highest price of AAPL stock on that date.
- **Low:** The lowest price of AAPL stock on that date.
- **Volume:** The trading volume of AAPL stock on that date.

The stock price data was aligned with the news data based on the publication date and time.

## 4.2 Sentiment-Stock Price Relationship

Once sentiment scores were calculated for each news headline and summary, the next step was to analyze whether sentiment correlated with stock price movements. For this, the following steps were followed:

- Calculate the **average sentiment** for news headlines and summaries on a daily basis.
- Compare sentiment scores with the **daily closing price** of AAPL.
- Investigate any patterns or correlations between sentiment and price fluctuations.

## 4.3 Data Visualization

Various plots were created to visualize the relationship between sentiment and stock prices. Key visualizations include:

- **Sentiment over time:** A line plot showing the daily sentiment scores for headlines and summaries.
- **Stock price over time:** A line plot showing the daily closing price of AAPL.
- **Sentiment vs. Stock Price:** Scatter plots to analyze if a strong correlation exists between sentiment scores and stock price changes.

## 5. Model Evaluation

### 5.1 Sentiment Analysis Evaluation

The effectiveness of sentiment analysis was evaluated by:

- **Accuracy of Sentiment Classification:** Checking if the polarity values correctly capture the overall tone of the news articles (positive, negative, or neutral).
- **Correlation with Stock Price Movements:** Analyzing the correlation between sentiment scores and stock price changes to determine if sentiment provides useful insights into market behavior.

### 5.2 Evaluation Metrics

Since this project is focused on exploratory data analysis rather than prediction, there are no formal metrics like accuracy, precision, or recall. Instead, the evaluation focused on:

- **Sentiment Score Distribution:** The distribution of sentiment scores for the headlines and summaries.
- **Correlation:** Pearson correlation coefficient between sentiment and stock price movements.

## 6. Results & Discussion

### 6.1 Sentiment Analysis Results

The sentiment analysis revealed that the headlines for Apple Inc. during January 2024 tended to be relatively neutral, with most sentiment scores falling between -0.2 and +0.2. However, certain spikes in positive or negative sentiment were observed, often linked to major news events (e.g., product launches, earnings reports, or regulatory news).

## 6.2 Sentiment and Stock Price Relationship

The analysis of sentiment against stock price changes showed:

- A **weak positive correlation** between headline sentiment and stock price movements. This suggests that positive news headlines may have a slight tendency to correlate with stock price increases.
- The **summary sentiment** exhibited a slightly stronger correlation with stock price, particularly when the sentiment was extremely positive or negative.

## 6.3 Limitations

While sentiment analysis provides interesting insights, the relationship between sentiment and stock price is not straightforward. Stock prices are influenced by various factors, such as market conditions, broader economic news, and investor behavior, which are not always captured by sentiment alone.

## 7. Conclusion

This project successfully demonstrated the use of sentiment analysis to analyze news related to Apple Inc. (AAPL) and its potential impact on stock price movements. The sentiment analysis revealed a general neutrality in the news, with occasional spikes in positive or negative sentiment. The correlation analysis between sentiment and stock price changes showed a weak positive relationship, suggesting that sentiment may have some impact on stock market behavior, but it is not the sole determining factor.

## 8.Original Code:

NLP (Natural Language Processing) for Stock Market News

```
*** Sentiment analysis ***
```

```
# Fetching Data Using Finnhub API
```

```
import finnhub
import pandas as pd
from textblob import TextBlob
```

```
# Setting-up client (using API key)
```

```
finnhub_client =
finnhub.Client(api_key='ctjss39r01quipmv1vr0ctjss39r01quipmv1vrg')
```

```
# Fetching news related to Apple (AAPL) from Finnhub API
```

```
res = finnhub_client.company_news('AAPL', _from="2024-01-01", to="2024-01-31")
```

```
# Converting the response into a pandas DataFrame for easier manipulation
```

```
z = pd.DataFrame(res)
```

```
# Adding sentiment scores for both headline and summary using TextBlob
```

```
z['headline_sentiment'] = z['headline'].apply(lambda x:
TextBlob(x).sentiment.polarity)
```

```
z['summary_sentiment'] = z['summary'].apply(lambda x:
TextBlob(x).sentiment.polarity)
```

```
# Displaying the first few rows of the DataFrame
```

```
print(z[['headline', 'summary', 'headline_sentiment',
'summary_sentiment']].head())
```

```
                                headline \
0  Tech earnings, jobless claims, mortgage rates:...
1  Apple to report Q1 earnings as investors focus...
2           12 Highest Quality Camera Phones in 2024
3  Big Tech earnings are here. Fasten your seat b...
4  Apple, Google Could Win the War for the Digita...
```

```
                                summary  headline_sentiment \
0  Yahoo Finance Live Co-Hosts Josh Lipton and Ju...      0.0
1  Apple will report its Q1 earnings after the be...      0.0
2  In this article, we will be going through the ...      0.0
3  It's earnings season, and Big Tech is taking a...      0.0
4  Big Tech wants your financial life on its apps...      0.4
```

```
summary_sentiment
```

```
0          0.265702
1          0.700000
2          0.240625
3          0.000000
4          0.050000
```

```
import pandas as pd
```

```
# Assuming 'z' is the DataFrame you have from Finnhub
# Converting 'datetime' column (epoch time) to a proper datetime format
z['date'] = pd.to_datetime(z['datetime'], unit='s').dt.date # Convert to date
```

```
# Calculating daily average sentiment for both headline and summary
daily_sentiment = z.groupby('date').agg({'headline_sentiment': 'mean',
'summary_sentiment': 'mean'}).reset_index()
```

```
# Displaying the result
print(daily_sentiment.head())
```

	date	headline_sentiment	summary_sentiment
0	2024-01-26	0.000000	0.500000
1	2024-01-27	0.157143	0.183368
2	2024-01-28	0.190949	0.097123
3	2024-01-29	0.088456	0.133244
4	2024-01-30	0.141400	0.140265

```
import matplotlib.pyplot as plt
```


```
# Plotting sentiment trends over time
plt.figure(figsize=(6, 4))
```

```
# Plotting both headline and summary sentiment
plt.plot(daily_sentiment['date'], daily_sentiment['headline_sentiment'],
label='Headline Sentiment', color='blue')
plt.plot(daily_sentiment['date'], daily_sentiment['summary_sentiment'],
label='Summary Sentiment', color='orange')
```

```
# Adding labels and title
plt.xlabel('Date')
plt.ylabel('Average Sentiment')
plt.title('Daily Sentiment Trend for Apple News')
plt.legend()
plt.xticks(rotation=45)
plt.grid(True)
```

```
plt.show()
```





```
import yfinance as yf
import pandas as pd
```

```
# Fetching historical data for AAPL (Apple)
stock_data = yf.download('AAPL', start='2024-01-01', end='2024-01-31',
progress=False)
```

```
# Displaying the stock data
print(stock_data.head())
```

Price Ticker Date	Close AAPL	High AAPL	Low AAPL	Open AAPL	Volume AAPL
2024-01-02	184.734970	187.521323	182.993502	186.237603	82488700
2024-01-03	183.351746	184.973804	182.535736	183.321893	58414500
2024-01-04	181.023163	182.197403	179.998185	181.261983	71983600
2024-01-05	180.296707	181.869006	179.291637	181.102771	62303300
2024-01-08	184.655365	184.695178	180.615161	181.202281	59144500

```
# Resetting the index, which will add the 'Date' column automatically
stock_data_reset = stock_data.reset_index()
```

```
# Verifying the structure of the DataFrame
print(stock_data_reset.head())
```

```
# Checking the closing price data
print(stock_data_reset[['Date', 'Close']].head())
```

Price Ticker	Date	Close AAPL	High AAPL	Low AAPL	Open AAPL	Volume AAPL
0	2024-01-02	184.734970	187.521323	182.993502	186.237603	82488700
1	2024-01-03	183.351746	184.973804	182.535736	183.321893	58414500
2	2024-01-04	181.023163	182.197403	179.998185	181.261983	71983600
3	2024-01-05	180.296707	181.869006	179.291637	181.102771	62303300
4	2024-01-08	184.655365	184.695178	180.615161	181.202281	59144500

Price Ticker	Date	Close AAPL
0	2024-01-02	184.734970
1	2024-01-03	183.351746
2	2024-01-04	181.023163
3	2024-01-05	180.296707
4	2024-01-08	184.655365

```
# Ensuring the 'date' column in daily_sentiment is in datetime format
daily_sentiment['date'] = pd.to_datetime(daily_sentiment['date'])
```

```
# Performing the merge after ensuring both dataframes have compatible columns
merged_df = pd.merge(daily_sentiment, stock_data_reset[['Date_',
'Close_AAPL']], left_on='date', right_on='Date_', how='inner')
```

```
# Displaying the merged data
print(merged_df.head())
```

	date	headline_sentiment	summary_sentiment	Date_	Close_AAPL
0	2024-01-26	0.000000	0.500000	2024-01-26	191.481903
1	2024-01-29	0.088456	0.133244	2024-01-29	190.795288
2	2024-01-30	0.141400	0.140265	2024-01-30	187.123260

```
import matplotlib.pyplot as plt
```

```
# Plotting sentiment trends and stock price
plt.figure(figsize=(6, 4))
```

```
# Plotting headline sentiment
plt.plot(merged_df['date'], merged_df['headline_sentiment'], label='Headline
Sentiment', color='blue', marker='o')
```

```
# Plotting summary sentiment
plt.plot(merged_df['date'], merged_df['summary_sentiment'], label='Summary
Sentiment', color='orange', marker='x')
```


```
# Plotting stock closing price
plt.plot(merged_df['date'], merged_df['Close_AAPL'], label='Stock Price',
color='green', linestyle='--')
```

```

# Adding labels and title
plt.xlabel('Date')
plt.ylabel('Sentiment / Stock Price')
plt.title('Sentiment vs Stock Price Over Time for Apple (January 2024)')
plt.legend(loc='upper left')
plt.xticks(rotation=45)
plt.grid(True)

plt.show()

```



```

# Calculating the correlation between sentiment scores and stock price

```

```

correlation_headline =
merged_df['headline_sentiment'].corr(merged_df['Close_AAPL'])
correlation_summary =
merged_df['summary_sentiment'].corr(merged_df['Close_AAPL'])

```

```

print(f"Correlation between headline sentiment and stock price:
{correlation_headline}")
print(f"Correlation between summary sentiment and stock price:
{correlation_summary}")

```

```

Correlation between headline sentiment and stock price: -0.867531692479687
Correlation between summary sentiment and stock price: 0.6082706222350976

```

```

from scipy.stats import pearsonr

# Calculating the correlation and p-value for headline sentiment vs stock
price
correlation_headline, p_value_headline =
pearsonr(merged_df['headline_sentiment'], merged_df['Close_AAPL'])

# Calculating the correlation and p-value for summary sentiment vs stock
price
correlation_summary, p_value_summary =
pearsonr(merged_df['summary_sentiment'], merged_df['Close_AAPL'])

# Output of results
print(f"Correlation between headline sentiment and stock price:
{correlation_headline}")
print(f"P-value for headline sentiment vs stock price: {p_value_headline}")

print()

print(f"Correlation between summary sentiment and stock price:
{correlation_summary}")
print(f"P-value for summary sentiment vs stock price: {p_value_summary}")

print()

# Interpretation of results
if p_value_headline < 0.05:
    print("Headline sentiment correlation is statistically significant.")
else:
    print("Headline sentiment correlation is not statistically significant.")

if p_value_summary < 0.05:
    print("Summary sentiment correlation is statistically significant.")
else:
    print("Summary sentiment correlation is not statistically significant.")

Correlation between headline sentiment and stock price: -0.8675316924796866
P-value for headline sentiment vs stock price: 0.33141043413211557

Correlation between summary sentiment and stock price: 0.6082706222350973
P-value for summary sentiment vs stock price: 0.5837270837256414

Headline sentiment correlation is not statistically significant.
Summary sentiment correlation is not statistically significant.

```

\*\*\* Impact of Sentiment on stock market price \*\*\*

```
import pandas as pd
import yfinance as yf

# Fetching stock data for Apple (AAPL)
stock_data = yf.download('AAPL', start='2024-01-01', end='2024-01-31',
progress=False)

# Checking the structure of stock_data (columns and index)
print("Stock data columns and index:")
print(stock_data.columns)
print(stock_data.index)

Stock data columns and index:
MultiIndex([( 'Close', 'AAPL'),
            ( 'High', 'AAPL'),
            (  'Low', 'AAPL'),
            ( 'Open', 'AAPL'),
            ('Volume', 'AAPL')],
            names=['Price', 'Ticker'])
DatetimeIndex(['2024-01-02', '2024-01-03', '2024-01-04', '2024-01-05',
               '2024-01-08', '2024-01-09', '2024-01-10', '2024-01-11',
               '2024-01-12', '2024-01-16', '2024-01-17', '2024-01-18',
               '2024-01-19', '2024-01-22', '2024-01-23', '2024-01-24',
               '2024-01-25', '2024-01-26', '2024-01-29', '2024-01-30'],
              dtype='datetime64[ns]', name='Date', freq=None)

import pandas as pd

# Flattening the multi-level columns
stock_data_reset = stock_data.copy()
stock_data_reset.columns = ['_'.join(col).strip() for col in
stock_data.columns]

# Checking the flattened columns
print(stock_data_reset.columns)

Index(['Close_AAPL', 'High_AAPL', 'Low_AAPL', 'Open_AAPL', 'Volume_AAPL'],
      dtype='object')

# Resetting the index to make 'Date' a regular column
stock_data_reset = stock_data_reset.reset_index()
```

*# Verifying the 'Date' column is now a regular column*

```
print(stock_data_reset.head())
```

	Date	Close_AAPL	High_AAPL	Low_AAPL	Open_AAPL	Volume_AAPL
0	2024-01-02	184.734970	187.521323	182.993502	186.237603	82488700
1	2024-01-03	183.351746	184.973804	182.535736	183.321893	58414500
2	2024-01-04	181.023163	182.197403	179.998185	181.261983	71983600
3	2024-01-05	180.296707	181.869006	179.291637	181.102771	62303300
4	2024-01-08	184.655365	184.695178	180.615161	181.202281	59144500

```
import matplotlib.pyplot as plt
```

*# Now plotting*

```
plt.figure(figsize=(6, 4))
```

*# Plotting headline sentiment*

```
plt.plot(merged_df['Date_'], merged_df['headline_sentiment'], label='Headline  
Sentiment', color='blue', marker='o')
```

*# Plotting summary sentiment*

```
plt.plot(merged_df['Date_'], merged_df['summary_sentiment'], label='Summary  
Sentiment', color='orange', marker='x')
```

*# Plotting stock close price (use 'Close\_AAPL' instead of 'close\_price')*

```
plt.plot(merged_df['Date_'], merged_df['Close_AAPL'], label='Stock Price',  
color='green', linestyle='--')
```

*# Adding labels and title*

```
plt.xlabel('Date')
```

```
plt.ylabel('Value')
```


```
plt.title('Sentiment vs Stock Price for Apple (January 2024)')
```

```
plt.legend()
```

```
plt.xticks(rotation=45)
```

```
plt.grid(True)
```

```
plt.show()
```



*# Calculating the correlation between sentiment scores and stock price*

```
correlation_headline =
merged_df['headline_sentiment'].corr(merged_df['Close_AAPL'])
correlation_summary =
merged_df['summary_sentiment'].corr(merged_df['Close_AAPL'])
```

```
print(f"Correlation between headline sentiment and stock price:
{correlation_headline}")
print(f"Correlation between summary sentiment and stock price:
{correlation_summary}")
```

```
Correlation between headline sentiment and stock price: -0.867531692479687
Correlation between summary sentiment and stock price: 0.6082706222350976
```

```
import scipy.stats as stats
```

*# Calculating Pearson correlation and p-value for headline sentiment vs stock price*

```
corr_headline, p_value_headline =
stats.pearsonr(merged_df['headline_sentiment'], merged_df['Close_AAPL'])
print(f"Correlation between headline sentiment and stock price:
{corr_headline}")
print(f"P-value for headline sentiment vs stock price: {p_value_headline}")
```

```
print()
```

```
# Calculating Pearson correlation and p-value for summary sentiment vs stock price
```

```
corr_summary, p_value_summary =  
stats.pearsonr(merged_df['summary_sentiment'], merged_df['Close_AAPL'])  
print(f"Correlation between summary sentiment and stock price:  
{corr_summary}")  
print(f"P-value for summary sentiment vs stock price: {p_value_summary}")  
  
print()
```

```
# Interpretation of the p-values
```

```
if p_value_headline < 0.05:  
    print("The correlation between headline sentiment and stock price is  
statistically significant.")  
else:  
    print("The correlation between headline sentiment and stock price is not  
statistically significant.")
```

```
if p_value_summary < 0.05:  
    print("The correlation between summary sentiment and stock price is  
statistically significant.")  
else:  
    print("The correlation between summary sentiment and stock price is not  
statistically significant.")
```

```
Correlation between headline sentiment and stock price: -0.8675316924796866  
P-value for headline sentiment vs stock price: 0.33141043413211557
```

```
Correlation between summary sentiment and stock price: 0.6082706222350973  
P-value for summary sentiment vs stock price: 0.5837270837256414
```

```
The correlation between headline sentiment and stock price is not  
statistically significant.
```

```
The correlation between summary sentiment and stock price is not  
statistically significant.
```

```
*** Time series prediction of stock value using various features ***
```

```
# Data Preparation and Preprocessing
```

```
import numpy as np  
import pandas as pd  
import yfinance as yf  
from sklearn.preprocessing import MinMaxScaler  
import matplotlib.pyplot as plt
```



```

# Fetching Stock Data
# Fetching historical data for Apple (AAPL) for the past 6 months
stock_data = yf.download('AAPL', start='2023-07-01', end='2024-01-31',
progress=False)

# Checking the stock data
print("Stock data head:")
print(stock_data.head())

# Fetching Sentiment Data (Ensure same Length as stock data)
# Generating random sentiment values for demonstration
dates = pd.date_range(start="2023-07-01", end="2024-01-31", freq='B')

sentiment_data = {
    'date': dates,
    'headline_sentiment': np.random.uniform(-1, 1, size=len(dates)),
    'summary_sentiment': np.random.uniform(-1, 1, size=len(dates)),
}

# Creating the sentiment dataframe
sentiment_df = pd.DataFrame(sentiment_data)
sentiment_df['date'] = pd.to_datetime(sentiment_df['date'])

# Merging sentiment data with stock data on date
stock_data_reset = stock_data.reset_index()[['Date', 'Close']]
stock_data_reset.columns = ['date', 'close_price']

merged_df = pd.merge(stock_data_reset, sentiment_df, on='date', how='inner')

# Checking the merged data
print("Merged data head:")
print(merged_df.head())

# Feature Engineering
# Adding a simple Moving Average (e.g., 5-day moving average) as additional
feature
merged_df['moving_avg_5'] = merged_df['close_price'].rolling(window=5).mean()

# Preparing the Data for LSTM Model
data = merged_df[['headline_sentiment', 'summary_sentiment', 'close_price',
'moving_avg_5']]

# Scaling the Features using MinMaxScaler
scaler = MinMaxScaler(feature_range=(0, 1))
scaled_data = scaler.fit_transform(data)

# Checking the scaled data shape
print(f"Scaled data shape: {scaled_data.shape}")

```

```
# Creating Sequences for LSTM (Using a Smaller Lookback Period)
```

```
lookback_period = 5 # Reduced Lookback period to 5 days
```

```
Stock data head:
```

Price	Close	High	Low	Open	Volume
Ticker	AAPL	AAPL	AAPL	AAPL	AAPL
Date					
2023-07-03	191.011810	192.421123	190.317065	192.321870	31458200
2023-07-05	189.890305	191.527883	189.185640	190.128504	46920300
2023-07-06	190.366684	190.575110	187.776323	188.411506	45094300
2023-07-07	189.245209	191.220240	188.808532	189.969727	46778000
2023-07-10	187.190781	188.560401	185.632587	187.835884	59922200

```
Merged data head:
```

	date	close_price	headline_sentiment	summary_sentiment
0	2023-07-03	191.011810	-0.283101	0.442617
1	2023-07-05	189.890305	0.587243	0.217056
2	2023-07-06	190.366684	-0.774843	0.793641
3	2023-07-07	189.245209	0.363401	-0.330838
4	2023-07-10	187.190781	-0.939702	0.794463

```
Scaled data shape: (146, 4)
```

```
# Model Definition, Training, and Evaluation
```

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import GRU, Dense, Dropout, Input
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint
```

```
X = []
```

```
y = []
```

```
# Creating sequences for training and prediction
```

```
for i in range(lookback_period, len(scaled_data)):
    X.append(scaled_data[i-lookback_period:i, :-1]) # Last 5 days of
sentiment and stock prices (excluding target)
    y.append(scaled_data[i, -1]) # Next day's closing price (target)
```

```
X, y = np.array(X), np.array(y)
```

```
# Reshaping X for LSTM (3D Shape: [samples, timesteps, features])
```

```
X = np.reshape(X, (X.shape[0], X.shape[1], X.shape[2]))
```

```
# Splitting the Data into Training and Testing Sets (80% training, 20%
testing)
```

```
train_size = int(len(X) * 0.8)
X_train, X_test = X[:train_size], X[train_size:]
y_train, y_test = y[:train_size], y[train_size:]
```

```
# Printing the shapes of training and testing data
```

```
print(f"Training data shape: {X_train.shape}")
```

```
print(f"Testing data shape: {X_test.shape}")
```

```

# Using GRU Model
model = Sequential()

# Adding Input Layer
model.add(Input(shape=(X_train.shape[1], X_train.shape[2])))

# Adding GRU Layer
model.add(GRU(units=50, return_sequences=True))
model.add(Dropout(0.2))

# Adding a second GRU Layer
model.add(GRU(units=50, return_sequences=False))
model.add(Dropout(0.2))

# Adding Output Layer (Single value for predicting the closing price)
model.add(Dense(units=1))

# Compiling the model
model.compile(optimizer='adam', loss='mean_squared_error')

# Defining Early Stopping and Model Checkpoint
early_stopping = EarlyStopping(monitor='val_loss', patience=10,
restore_best_weights=True)
model_checkpoint = ModelCheckpoint('best_model.keras', monitor='val_loss',
save_best_only=True)

# Training the Model with Early Stopping and Model Checkpointing
model.fit(X_train, y_train, epochs=50, batch_size=32, validation_split=0.2,
callbacks=[early_stopping, model_checkpoint])

```

#### # Evaluating the Model

```

test_loss = model.evaluate(X_test, y_test)
print(f'Test Loss: {test_loss}')

```

Training data shape: (112, 5, 3)

Testing data shape: (29, 5, 3)

Epoch 1/50

3/3 ————— 2s 110ms/step - loss: 0.1784 -  
val\_loss: 0.2727

Epoch 2/50

3/3 ————— 0s 16ms/step - loss: 0.0855 -  
val\_loss: 0.0812

Epoch 3/50

3/3 ————— 0s 17ms/step - loss: 0.0643 -  
val\_loss: 0.0251

Epoch 4/50

3/3 ————— 0s 19ms/step - loss: 0.0624 -  
val\_loss: 0.0240

Epoch 5/50

3/3 ————— 0s 9ms/step - loss: 0.0493 -  
val\_loss: 0.0442

Epoch 6/50

3/3 ————— 0s 9ms/step - loss: 0.0301 -  
val\_loss: 0.0651  
Epoch 7/50  
3/3 ————— 0s 9ms/step - loss: 0.0349 -  
val\_loss: 0.0709  
Epoch 8/50  
3/3 ————— 0s 9ms/step - loss: 0.0266 -  
val\_loss: 0.0592  
Epoch 9/50  
3/3 ————— 0s 9ms/step - loss: 0.0229 -  
val\_loss: 0.0355  
Epoch 10/50  
3/3 ————— 0s 17ms/step - loss: 0.0137 -  
val\_loss: 0.0151  
Epoch 11/50  
3/3 ————— 0s 17ms/step - loss: 0.0097 -  
val\_loss: 0.0047  
Epoch 12/50  
3/3 ————— 0s 17ms/step - loss: 0.0071 -  
val\_loss: 0.0017  
Epoch 13/50  
3/3 ————— 0s 9ms/step - loss: 0.0090 -  
val\_loss: 0.0021  
Epoch 14/50  
3/3 ————— 0s 9ms/step - loss: 0.0081 -  
val\_loss: 0.0039  
Epoch 15/50  
3/3 ————— 0s 13ms/step - loss: 0.0058 -  
val\_loss: 0.0046  
Epoch 16/50  
3/3 ————— 0s 12ms/step - loss: 0.0059 -  
val\_loss: 0.0018  
Epoch 17/50  
3/3 ————— 0s 17ms/step - loss: 0.0068 -  
val\_loss: 1.9410e-04  
Epoch 18/50  
3/3 ————— 0s 18ms/step - loss: 0.0087 -  
val\_loss: 1.6668e-04  
Epoch 19/50  
3/3 ————— 0s 9ms/step - loss: 0.0054 -  
val\_loss: 3.1594e-04  
Epoch 20/50  
3/3 ————— 0s 9ms/step - loss: 0.0069 -  
val\_loss: 7.9997e-04  
Epoch 21/50  
3/3 ————— 0s 9ms/step - loss: 0.0054 -  
val\_loss: 0.0024  
Epoch 22/50  
3/3 ————— 0s 9ms/step - loss: 0.0045 -  
val\_loss: 0.0039  
Epoch 23/50

```

3/3 _____ 0s 9ms/step - loss: 0.0066 -
val_loss: 0.0033
Epoch 24/50
3/3 _____ 0s 9ms/step - loss: 0.0044 -
val_loss: 0.0031
Epoch 25/50
3/3 _____ 0s 9ms/step - loss: 0.0059 -
val_loss: 0.0025
Epoch 26/50
3/3 _____ 0s 9ms/step - loss: 0.0054 -
val_loss: 0.0021
Epoch 27/50
3/3 _____ 0s 10ms/step - loss: 0.0052 -
val_loss: 0.0029
Epoch 28/50
3/3 _____ 0s 11ms/step - loss: 0.0054 -
val_loss: 0.0024
1/1 _____ 0s 13ms/step - loss: 0.0018
Test Loss: 0.0018108411459252238

```

#### *# Model Evaluation and Results Visualization*

```
from sklearn.metrics import r2_score, mean_squared_error
```

#### *# Making Predictions on the Test Set*

```
y_pred = model.predict(X_test)
```

#### *# Inverse Transforming the Predictions and Actuals*

```

y_pred_rescaled = scaler.inverse_transform(np.column_stack((y_pred, y_pred,
y_pred, y_pred)))[:, 0]
y_test_rescaled = scaler.inverse_transform(np.column_stack((y_test, y_test,
y_test, y_test)))[:, 0]

```

#### *# Calculating RMSE and R-squared*

```

rmse = np.sqrt(mean_squared_error(y_test_rescaled, y_pred_rescaled))
r2 = r2_score(y_test_rescaled, y_pred_rescaled)

```

```
print(f'Root Mean Squared Error (RMSE): {rmse}')
```

```
print(f'R-squared: {r2}')
```

#### *# Calculating Accuracy Based on Direction (Up or Down)*

```

y_test_direction = np.sign(np.diff(y_test_rescaled)) # Direction for actual
values
y_pred_direction = np.sign(np.diff(y_pred_rescaled)) # Direction for
predicted values

```

#### *# Ensuring both directions have the same length by trimming the first element*

```

y_test_direction = y_test_direction[1:]
y_pred_direction = y_pred_direction[1:]

```

#### *# Calculating accuracy based on direction prediction*

```
correct_direction = np.sum(y_test_direction == y_pred_direction)
```

```

accuracy = correct_direction / len(y_test_direction)
print(f'Accuracy based on direction: {accuracy * 100:.2f}%')

# Plotting the Results
plt.figure(figsize=(6, 4))


# Plotting actual stock price
plt.plot(y_test_rescaled, color='red', label='Actual Stock Price')

# Plotting predicted stock price
plt.plot(y_pred_rescaled, color='blue', label='Predicted Stock Price')

plt.title('Stock Price Prediction using GRU')
plt.xlabel('Time')
plt.ylabel('Stock Price')
plt.legend()
plt.show()

```

1/1 ————— 0s 160ms/step  
 Root Mean Squared Error (RMSE): 0.08410660097012336  
 R-squared: 0.9313721918943583  
 Accuracy based on direction: 96.30%



**\*\* Named Entity Recognition on news data \*\***

```

import pandas as pd
import spacy

# Loading the pre-trained spaCy NER model
nlp = spacy.load("en_core_web_sm")

data = {
    'date': ['2023-07-03', '2023-07-05', '2023-07-06'],
    'headline': [
        'Apple launches new iPhone in the US market',
        'Tesla stocks rise as Elon Musk hints at new product',
        'Google to invest heavily in AI research in Europe'
    ]
}

# Creating the DataFrame
news_df = pd.DataFrame(data)

# Initializing a list to store the extracted entities
entities_list = []

# Looping through each headline to extract named entities
for index, row in news_df.iterrows():
    headline = row['headline']

    # Processing the headline text using spaCy NER
    doc = nlp(headline)
    entities = {'ORG': [], 'PERSON': [], 'GPE': [], 'MONEY': [], 'PRODUCT':
[], 'LOC': []}

    # Extracting entities from the headline
    for ent in doc.ents:
        if ent.label_ == 'ORG':
            entities['ORG'].append(ent.text)
        elif ent.label_ == 'PERSON':
            entities['PERSON'].append(ent.text)
        elif ent.label_ == 'GPE':
            entities['GPE'].append(ent.text)
        elif ent.label_ == 'MONEY':
            entities['MONEY'].append(ent.text)
        elif ent.label_ == 'PRODUCT':
            entities['PRODUCT'].append(ent.text)
        elif ent.label_ == 'LOC':
            entities['LOC'].append(ent.text)

    # Appending the entities to the list
    entities_list.append(entities)

# Adding the extracted entities to the DataFrame
news_df['entities'] = entities_list

```

```
# Displaying the updated DataFrame with the extracted entities
```

```
print(news_df[['date', 'headline', 'entities']])
```

```
      date      headline \
0  2023-07-03  Apple launches new iPhone in the US market
1  2023-07-05  Tesla stocks rise as Elon Musk hints at new pr...
2  2023-07-06  Google to invest heavily in AI research in Europe
```

```
      entities
0  {'ORG': ['Apple'], 'PERSON': [], 'GPE': ['US']}...
1  {'ORG': [], 'PERSON': ['Elon Musk'], 'GPE': []}...
2  {'ORG': [], 'PERSON': [], 'GPE': ['AI'], 'MONE'...
```

```
# Extracting and count occurrences of entities
```

```
org_entities = [ent for sublist in news_df['entities'] for ent in
sublist['ORG']]
person_entities = [ent for sublist in news_df['entities'] for ent in
sublist['PERSON']]
product_entities = [ent for sublist in news_df['entities'] for ent in
sublist['PRODUCT']]
gpe_entities = [ent for sublist in news_df['entities'] for ent in
sublist['GPE']]
```

```
# Frequency counts
```

```
org_count = pd.Series(org_entities).value_counts()
person_count = pd.Series(person_entities).value_counts()
product_count = pd.Series(product_entities).value_counts()
gpe_count = pd.Series(gpe_entities).value_counts()
```

```
# Displaying counts
```

```
print("Top Organizations:\n", org_count.head())
print("Top People:\n", person_count.head())
print("Top Products:\n", product_count.head())
```

```
Top Organizations:
```

```
Apple    1
```

```
Name: count, dtype: int64
```

```
Top People:
```

```
Elon Musk    1
```

```
Name: count, dtype: int64
```

```
Top Products:
```

```
Series([], Name: count, dtype: int64)
```

```
from textblob import TextBlob
```

```
# Calculating sentiment of each headline
```

```
news_df['headline_sentiment'] = news_df['headline'].apply(lambda x:
TextBlob(x).sentiment.polarity)
```



```
# Filtering the sentiment of articles mentioning a specific organization
apple_news = news_df[news_df['entities'].apply(lambda x: 'Apple' in
x['ORG'])]
apple_sentiment = apple_news['headline_sentiment'].mean()
```

```
print(f"Average sentiment for Apple: {apple_sentiment}")
```

Average sentiment for Apple: 0.13636363636363635

```
import matplotlib.pyplot as plt
```

```
# Assuming that `news_df` contains a column for dates and sentiment
news_df['date'] = pd.to_datetime(news_df['date']) # Ensure dates are in
datetime format
```

```
# Assuming we have a 'headline_sentiment' column for sentiment scores
news_df['headline_sentiment'] = news_df['headline'].apply(lambda x:
TextBlob(x).sentiment.polarity)
```

```
# Plotting sentiment over time
plt.figure(figsize=(6, 4))
plt.plot(news_df['date'], news_df['headline_sentiment'], label='Headline
Sentiment', marker='o', color='blue')
plt.xlabel('Date')
plt.ylabel('Sentiment Score')
plt.title('Apple Sentiment Over Time')
plt.xticks(rotation=45)
plt.grid(True)
plt.legend()
plt.show()
```

